

Advanced Database

Phase 2

JANUARY 16

Ahmad Nader Adel – 1162296

Kamel Mohsen Kamel – 1162325

Ahmed Mohamed Amrawy – 1162173

Fatema Othman Mahmoud – 1162040



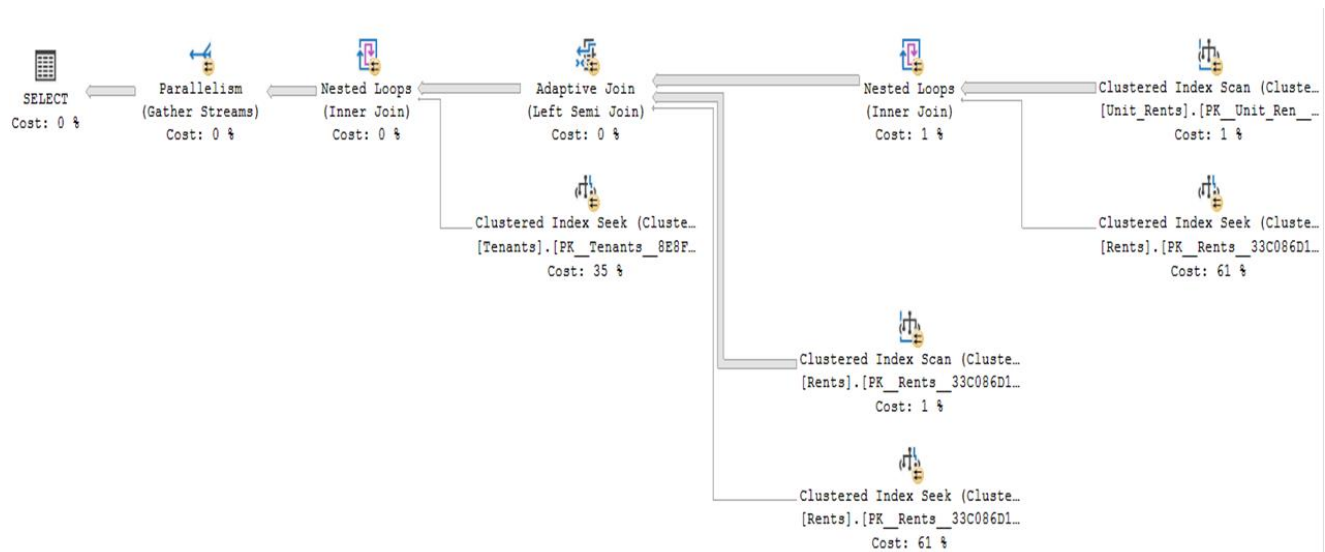
Contents

Query 1	3
Before Optimization:	3
After Optimization	4
Optimization Done:	5
Query 2	6
Before Optimization:	6
After Optimization	7
Optimization Done:	8
Query 3	9
Optimization Done:	10
New Database Statistics	11
Time Analysis	12
Time Analysis	13
Optimizations	14
The Enhancement in the Schema	14
The Enhancement of Memory Management	14
The Modification in Indexes	14
Conclusion	15

Query 1

Before Optimization:

```
SELECT T.Tenant_Id,T.Tenant_FirstName,T.Tenant_LastName
FROM Unit_Rents UR
INNER LOOP JOIN Rents R ON
UR.Unit_Rents_RentId = R.Rent_Id
INNER LOOP JOIN Tenants T ON
UR.Unit_Rents_TenantId = T.Tenant_Id
WHERE R.Rent_DueDate >= '1900-01-01'
AND R.Rent_DueDate <= '2020-12-31'
    AND R.Rent_Id in
    (
        SELECT R.Rent_Id
        FROM Rents R
        WHERE R.Rent_Statues = 'False'
    )
```

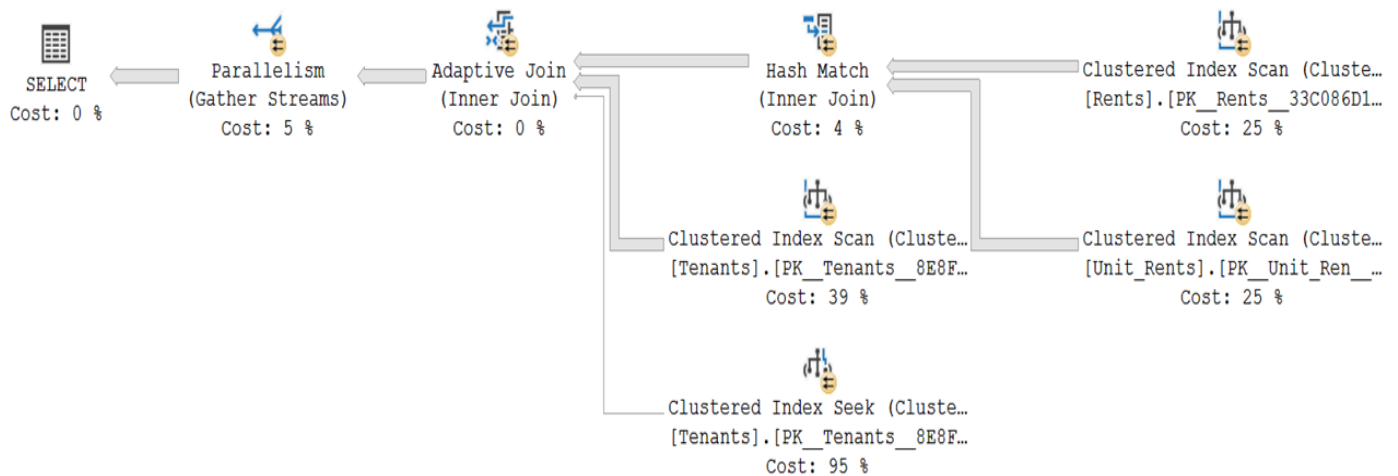


Rows Affected = 980,165 rows

Query 1

After Optimization

```
SELECT T.Tenant_Id,T.Tenant_FirstName,T.Tenant_LastName
FROM Unit_Rents UR
INNER JOIN Rents R ON
UR.Unit_Rents_RentId = R.Rent_Id
INNER JOIN Tenants T ON
UR.Unit_Rents_TenantId = T.Tenant_Id
WHERE R.Rent_DueDate >= '1900-01-01'
AND R.Rent_DueDate <= '2020-12-31'
AND R.Rent_Statues = 'False'
```



Query 1

Optimization Done:

- Inner Join Instead of Nested loop join
- Non Clustered Index on : **Rent_DueDate**
- Non Clustered Index on : **Rent_Statues**

<u>Action</u>	<u>CPU TIME</u>	<u>ELAPSED TIME</u>
<u>Query 1(No optimization)</u>	<u>9468 ms</u>	<u>8341 ms</u>
<u>Query 1(Inner Join and rewriting the query)</u>	<u>2667 ms</u>	<u>7026 ms</u>
<u>Query 1(Indexes Added)</u>	<u>2216 ms</u>	<u>6920 ms</u>
<u>Query 1(Mongo)</u>	<u>144,301 ms</u>	

Parallel Query Processing Report (theoretically):

- We can make parallel execution as follows:
 - Clustered index scan on rents table
 - Clustered index scan on unit_rents table
 - Clustered index scan on tenants
 - Clustered index seek on tenants table

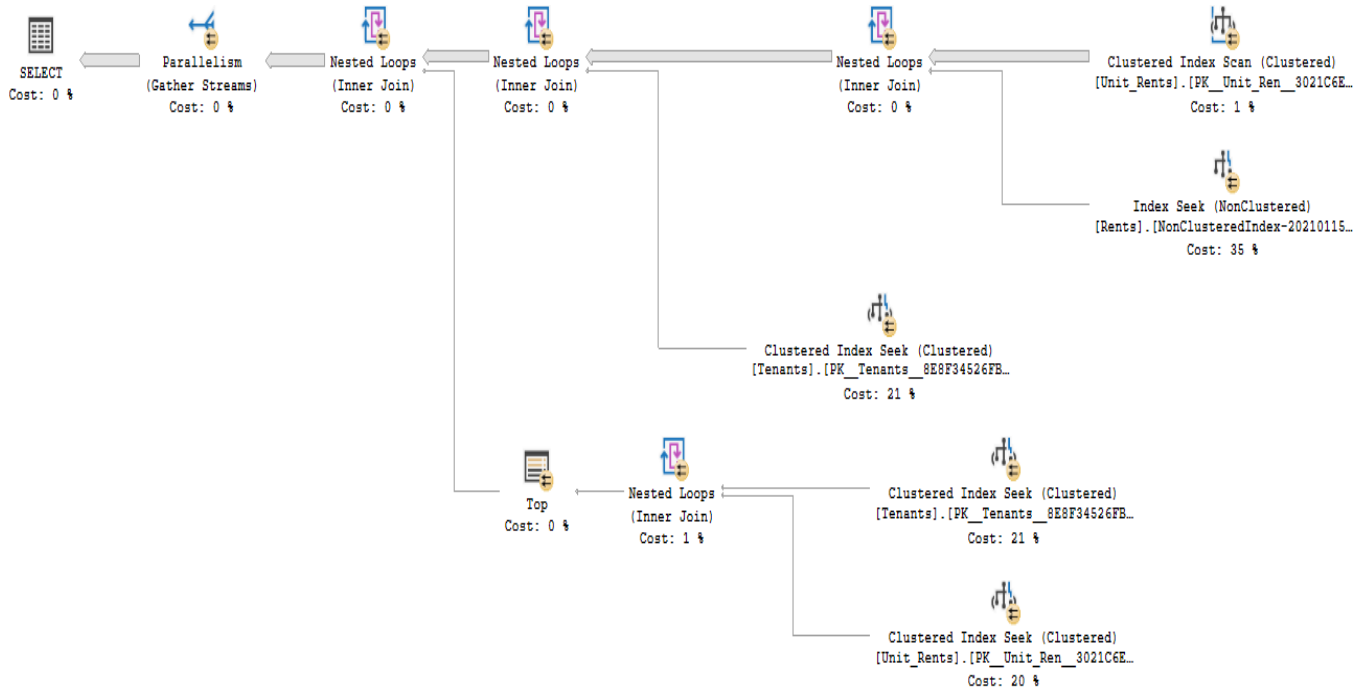
Query 2

Before Optimization:

```

SELECT T.Tenant_FirstName, T.Tenant_LastName
FROM Tenants T, Rents R , Unit_Rents UR
WHERE R.Rent_Statues = 'False'
      AND UR.Unit_Rents_TenantId = T.Tenant_Id
      AND UR.Unit_Rents_RentId = R.Rent_Id
      AND T.Tenant_Id
      in (
          SELECT      T.Tenant_Id
          FROM Unit_Rents UR, Units U,Tenants T
          WHERE T.Tenant_Id = UR.Unit_Rents_TenantId
                AND U.Unit_Id = UR.Unit_Rents_UnitId
      )

```

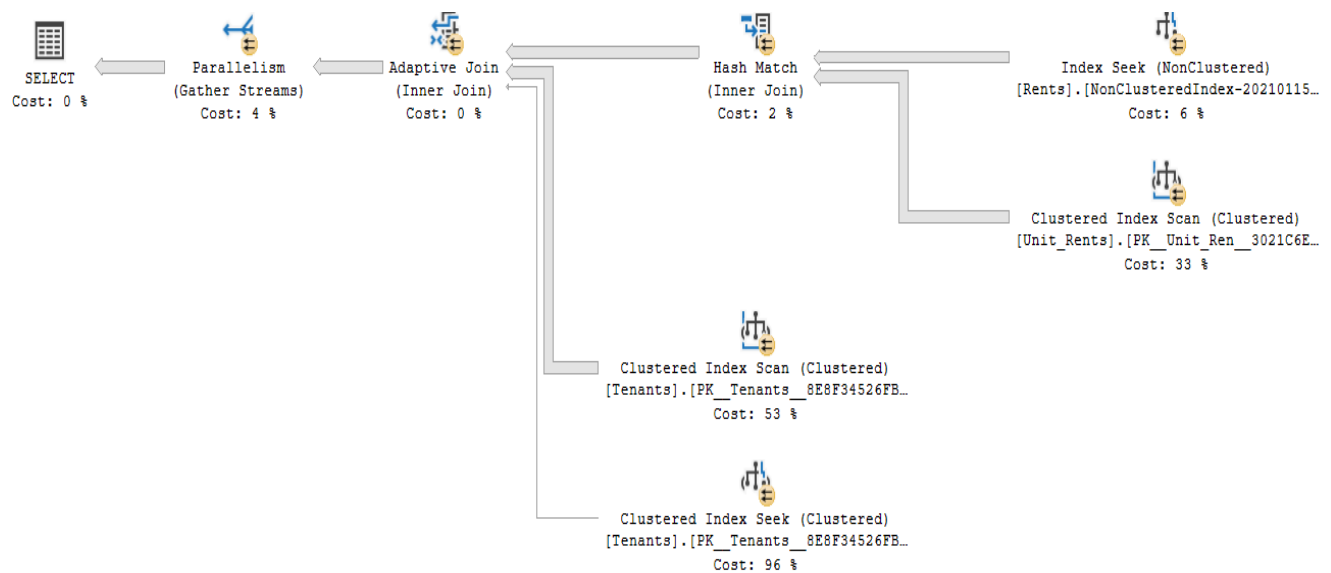


Rows Affected = 990,071 rows

Query 2

After Optimization

```
SELECT T.Tenant_FirstName, T.Tenant_LastName
FROM
Unit_Rents UR inner join Rents R on UR.Unit_Rents_RentId = R.Rent_Id
inner join
Units U on U.Unit_Id = UR.Unit_Rents_UnitId
inner join
Tenants T on T.Tenant_Id = UR.Unit_Rents_TenantId
WHERE R.Rent_Statues = 'False'
```



Query 2

Optimization Done:

- Inner Join Instead of Nested loop join
- Non Clustered Index on : **Rent_Statues**

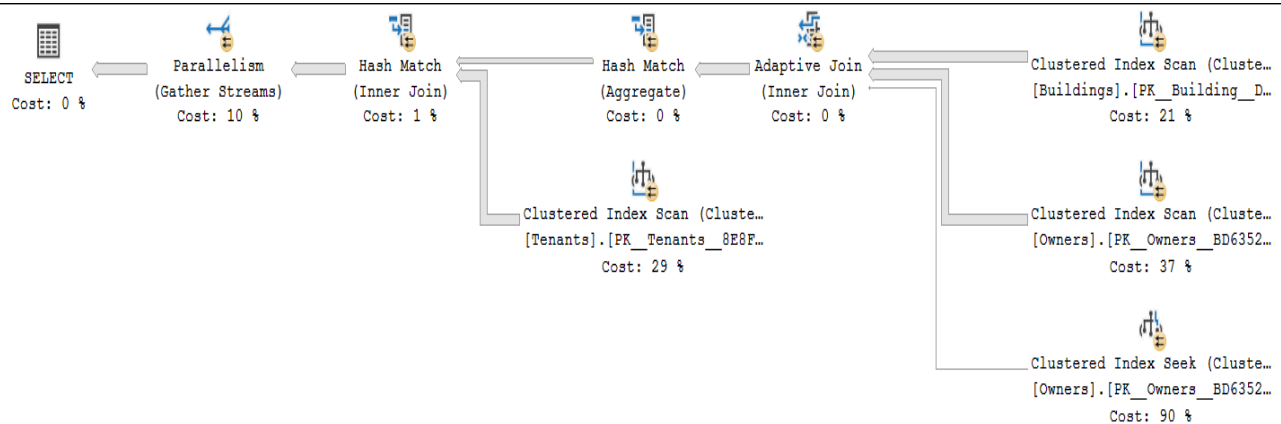
<u>Action</u>	<u>CPU TIME</u>	<u>ELAPSED TIME</u>
<u>Query 2(No optimization)</u>	<u>9140 ms</u>	<u>5608 ms</u>
<u>Query 2(Inner Join and rewriting the query)</u>	<u>1705 ms</u>	<u>7084 ms</u>
<u>Query 2(Indexes Added)</u>	<u>1372 ms</u>	<u>5463 ms</u>
<u>Query 2(Mongo)</u>	<u>99,822 ms</u>	

Parallel Query Processing Report (theoretically):

- We can make parallel execution as follows:
 - Non Clustered index seek on rents table
 - Clustered index scan on unit_rents table
 - Clustered index scan on tenants
 - Clustered index seek on tenants table

Query 3

```
SELECT *
FROM Tenants T
WHERE T.Tenant_FirstName in
(
    SELECT O.Owner_FirstName
    From Buildings B, Owners O
    WHERE B.Building_OwnerId = O.Owner_Id
)
```



Rows Affected = 1,979,936 rows

Query 3

Optimization Done:

- Changed **Owner_FirstName**, and **Tenant_FirstName** from VARCHAR(20) to VARCHAR(10).
- Added the query as a stored procedure

<u>Action</u>	<u>CPU TIME</u>	<u>ELAPSED TIME</u>
<u>Query 3(No optimization)</u>	<u>4907 ms</u>	<u>22588 ms</u>
<u>Query 3(VARCHAR(10))</u>	<u>4836 ms</u>	<u>16785 ms</u>
<u>Query 3(stored procedure)</u>	<u>5047 ms</u>	<u>15902 ms</u>
<u>Query 1(Mongo)</u>	<u>310,585 ms</u>	

Parallel Query Processing Report (theoretically):

- We can make parallel execution as follows:


- Clustered index scan on buildings table
- Clustered index scan on Owners table
- Clustered index seek on Owners table
- Clustered index scan on Tenants table

New Database Statistics

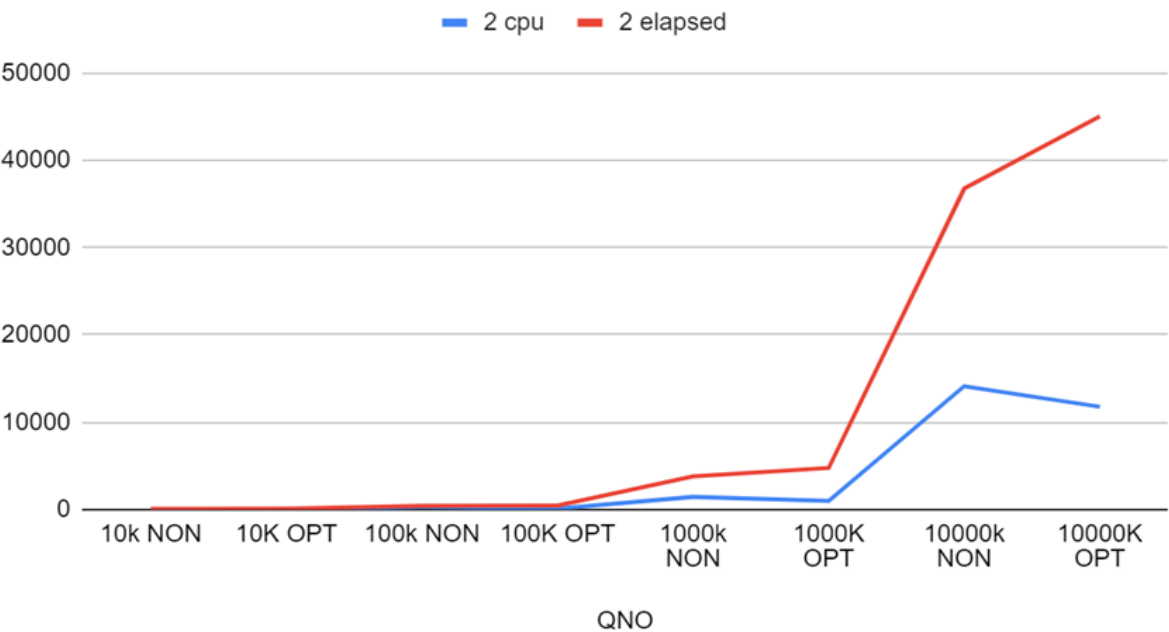
- The Tenants Max row size decreased by 10 bytes
- The Owners Max row size decreased by 10 bytes
- 2 indexes were added to the rents table

<u>Table Name</u>	<u>Row Count</u>	<u>Main Key</u>	<u>Indexes</u>	<u>FK</u>	<u>Identity Column</u>	<u>Max Row Size (Bytes)</u>
<u>Buildings</u>	<u>1076560</u>	<u>YES</u>	<u>2</u>	<u>1</u>	<u>YES</u>	<u>96</u>
<u>Contacts</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>0</u>	<u>YES</u>	<u>134</u>
<u>Employees</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>3</u>	<u>YES</u>	<u>139</u>
<u>Expenses</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>1</u>	<u>YES</u>	<u>43</u>
<u>Leases</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>1</u>	<u>YES</u>	<u>46</u>
<u>Owners</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>0</u>	<u>YES</u>	<u>86</u>
<u>Rents</u>	<u>2000000</u>	<u>YES</u>	<u>3</u>	<u>2</u>	<u>YES</u>	<u>24</u>
<u>Tenants</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>1</u>	<u>YES</u>	<u>62</u>
<u>Unit Rents</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>3</u>	<u>YES</u>	<u>12</u>
<u>Units</u>	<u>2000000</u>	<u>YES</u>	<u>1</u>	<u>1</u>	<u>YES</u>	<u>65</u>
<u>Users</u>	<u>969</u>	<u>YES</u>	<u>1</u>	<u>0</u>	<u>YES</u>	<u>18</u>

Time Analysis

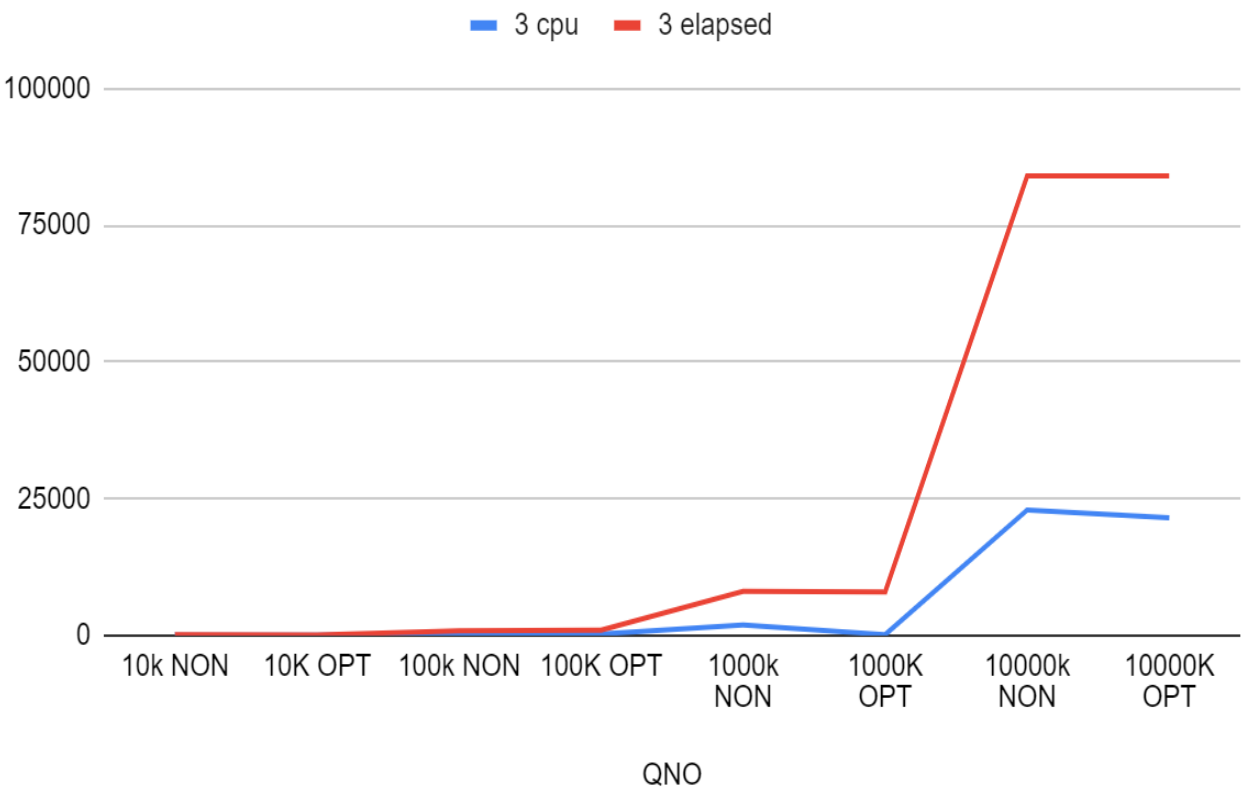


Query 2 CPU & Elapsed



Time Analysis

Query 3 CPU & Elapsed



Optimizations

The Enhancement in the Schema

- No modifications in the schema since all the relations between the tables normalized before

The Enhancement of Memory Management

- Added some queries as stored procedure to enhance the memory and cache usage.
- We have changed the type/reduced the size of the some column like:
 - Changed Owner_FirstName, and Tenant_FirstName from CHAR(20) to VARCHAR(10).

The Modification in Indexes

- We have added some indexes to boost the selection and update of data
 - Non Clustered Index on : **Rent_DueDate**
 - Non Clustered Index on : **Rent_Statues**
 - Non Clustered Index on : **Rent_Statues**

Conclusion

- In time Optimization: index tuning and query optimization affected the time greatly, it improved by almost 50% in most of the queries
- Simple queries took a lot of time even though the schema was normalized correctly, the main key was which columns to choose as the indexes, specially from foreign keys
- Memory Enhancements were the same in most of the queries like data selection and data update, but in data selection, the most complicated query, there was a great improvement and reduction in memory usage.
- In data insertion query, as the time optimization had great results, memory usage was notably increasing in change.
- When trying different database sizes, there was barely any increase in memory or time, which is expected due to all the made optimizations.
- The overall optimization had good results