



**CMPN-401**

**Advanced Database  
Systems**



**Cairo University**

**Faculty of Engineering**

# ADB Project report

**Submitted by:**

- Omar Mohamed Ahmed Abdellatif 1162003
- Mahmoud Atwa 1162014
- Ahmed Mohamed Ali Gad Hashish 1162017
- Ahmed Mohamed Ahmed Hassan 1162057

**Submitted to:**

- Dr. Michael Nawar
- Eng. Mostafa Mahmoud

## Contents

Specs of the used device:.....	2
Before and After Optimization Query Statistics: .....	3
Query 1:.....	3
Before Optimization:.....	3
After Optimization: .....	4
Query 2:.....	5
Before Optimization:.....	5
After Optimization: .....	5
Query 3:.....	7
Before Optimization:.....	7
After Optimization: .....	7
Optimization Details: .....	9
The new database statistics after modification:.....	9
• Disk Usage by Partition: .....	9
• Disk Usage by Table: .....	9
The enhancements made: .....	10
The enhancements in the schema: .....	10
The enhancement in memory management: .....	10
The modification in the indexes:.....	10
Modifications on the query statements: .....	10
Validation Details .....	11
Effect of modifications on query performance:.....	11
Effect of database size on query performance: .....	11
SQL vs NoSQL: .....	12
Equivalent NoSQL queries used: .....	12
Comments and Recommendations:.....	13

## Specs of the used device:

- Processor: Intel Core i7-9750H CPU @ 2.60GHz
- RAM: 16.0 GB
- OS: Windows 10 64-bit

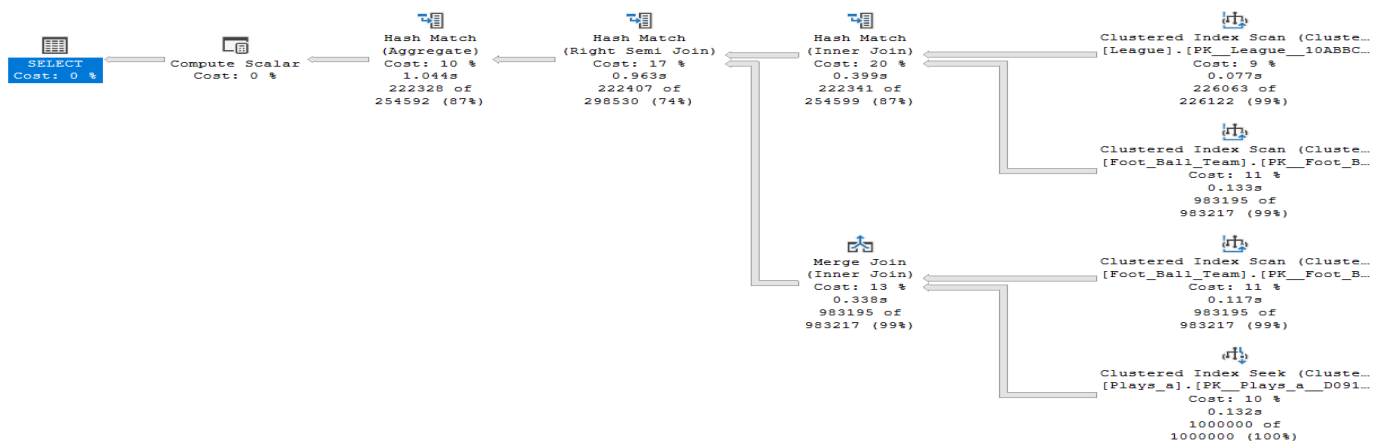
# Before and After Optimization Query Statistics:

## Query 1:

### Before Optimization:

```
SELECT COUNT(*)
FROM dbo.Plays_a, dbo.Foot_Ball_Team, dbo.Game
WHERE dbo.Plays_a.TeamID = dbo.Foot_Ball_Team.TeamID AND dbo.Plays_a.GameID = dbo.Game.GameID
AND Foot_Ball_Team.Goals IN
(SELECT dbo.Foot_Ball_Team.Goals
FROM dbo.Foot_Ball_Team, dbo.League
WHERE dbo.Foot_Ball_Team.LeagueID = dbo.League.LeagueID AND dbo.League.year > 2007
)
and Foot_Ball_Team.Goals > 15000000
GROUP BY dbo.Foot_Ball_Team.Goals
```

### • Execution Plan:



### • Theoretical parallel query processing:

The following can be executed in parallel:

- Clustered Index scan on League table
- Two clustered Index scans on Foot\_Ball\_Team table
- Clustered Index seek on Play\_a table

Also the following can be executed in parallel:

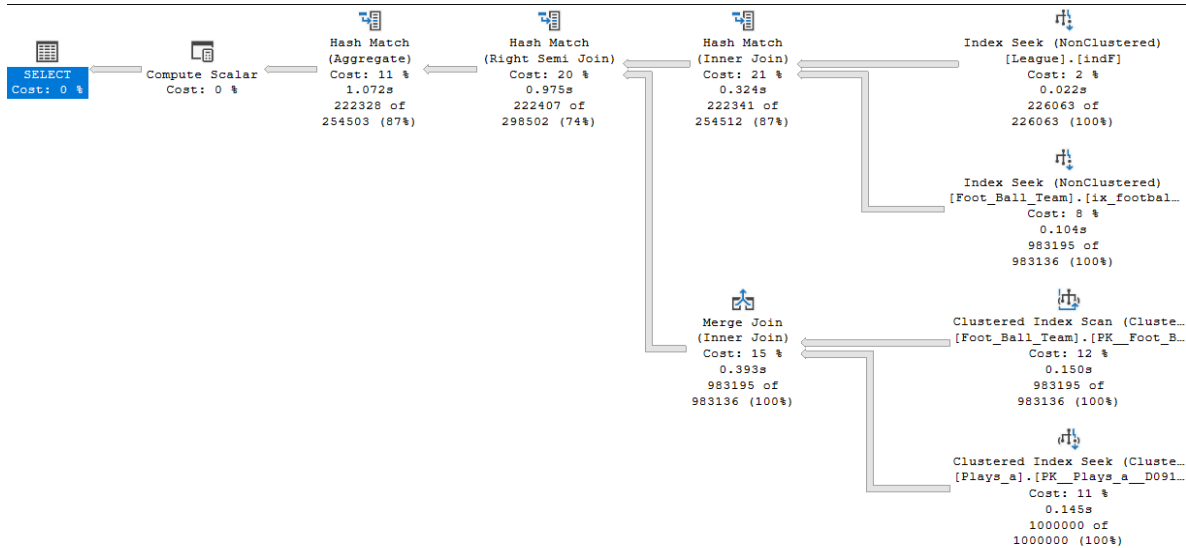
- Hash Match
- Merge Join

## After Optimization:

- Optimization using Indexing:

```
CREATE INDEX indf ON dbo.League(year)
CREATE INDEX ix_footballA ON dbo.Foot_Ball_Team(Goals,LeagueID)
```

- Execution plan:



- Theoretical parallel query processing:

The following can be executed in parallel:

- Index seek on League table
- Index seek on Foot\_ball\_Team table
- Clustered Index scan on Foot\_ball\_Team table
- Clustered Index seek on Plays\_a table

Also the following can be executed in parallel:

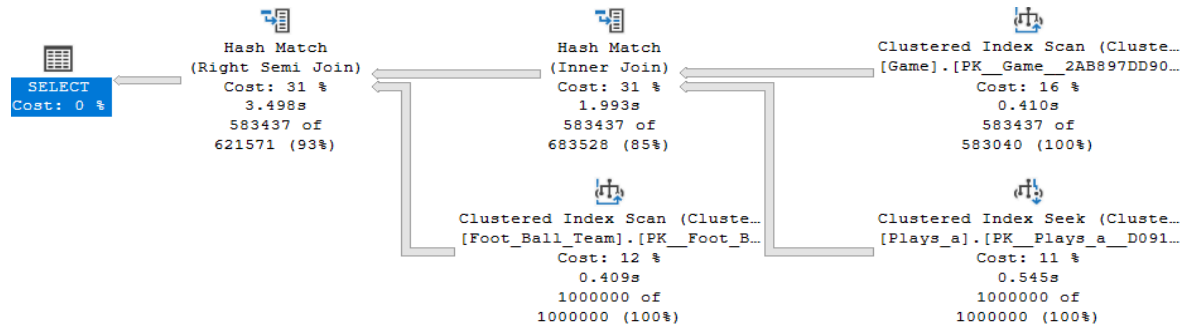
- Hash Match
- Merge Join

## Query 2:

### Before Optimization:

```
SELECT dbo.Foot_Ball_Team.Name
FROM dbo.Foot_Ball_Team
WHERE TeamID IN (SELECT dbo.Plays_a.TeamID
FROM dbo.Plays_a, dbo.Game
WHERE dbo.Plays_a.GameID = dbo.Game.GameID AND dbo.Game.Monthhh > 5)
```

- Execution Plan:



- Theoretical parallel query processing:

The following can be executed in parallel:

- Clustered Index scan on Game table
- Clustered Index seek on Plays\_a table
- Clustered Index scan on Foot\_Ball\_Team table

### After Optimization:

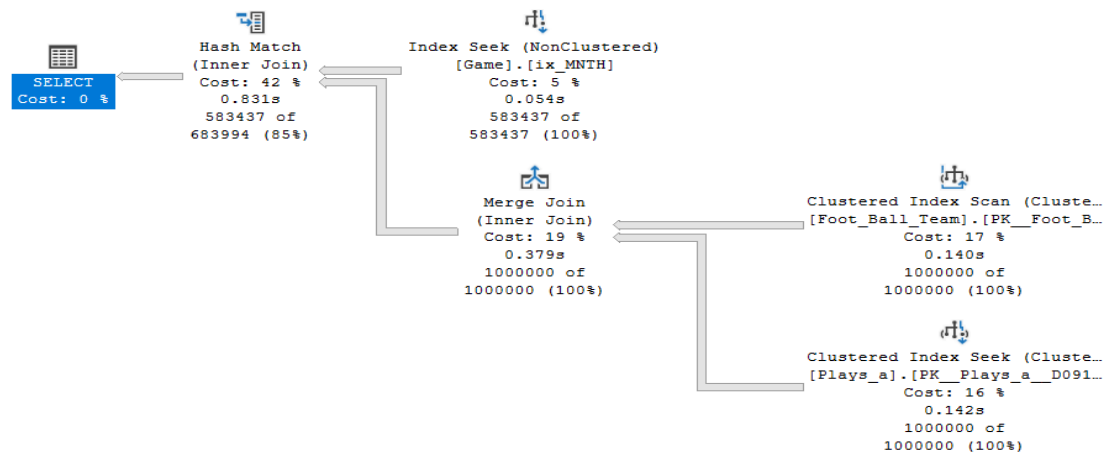
- Optimization using Indexing:

```
CREATE INDEX ix_MNTH ON dbo.Game(Monthhh)
```

- Optimization using query re-writing:

```
SELECT dbo.Foot_Ball_Team.Name
FROM dbo.Foot_Ball_Team, dbo.Plays_a, dbo.Game
WHERE Foot_Ball_Team.TeamID = Plays_a.TeamID AND Plays_a.GameID = Game.GameID AND
dbo.Game.Monthhh > 5
```

- Execution plan:



- Theoretical parallel query processing:

The following can be executed in parallel:

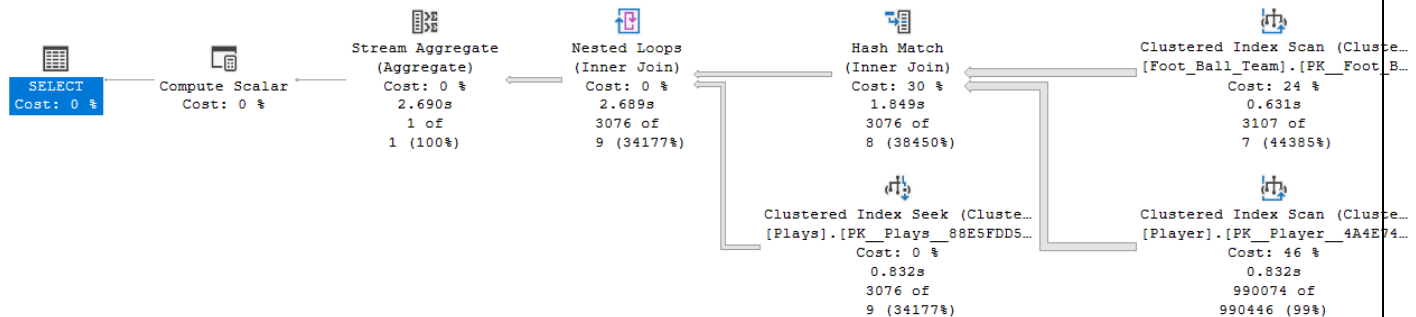
- Clustered Index scan on Foot\_Ball\_Team table
- Clustered Index seek on Plays\_a table
- Index seek on Game table

# Query 3:

## Before Optimization:

```
SELECT COUNT(*)
FROM dbo.Plays
INNER JOIN dbo.Player ON Player.PlayerID = Plays.PlayerID
INNER JOIN dbo.Foot_Ball_Team ON Foot_Ball_Team.TeamID = Player.TeamID
WHERE SUBSTRING(dbo.Foot_Ball_Team.Name,1,2) = 'AB'
```

- Execution Plan:



- Theoretical parallel query processing:

The following can be executed in parallel:

- Clustered Index scan on Foot\_Ball\_Team table
- Clustered Index scan on Player table
- Clustered Index seek on Plays table

## After Optimization:

- Optimization using Indexing:

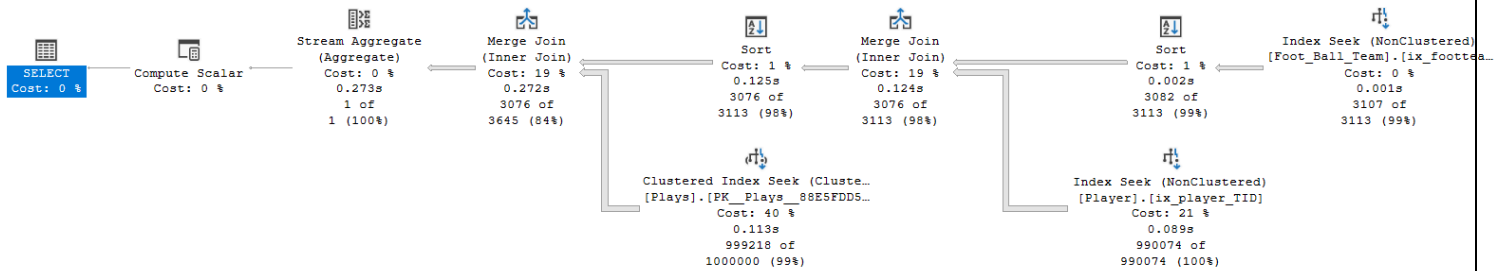
```
CREATE INDEX ix_player_TID ON dbo.Player(TeamID)
CREATE INDEX ix_footteam_team_name ON dbo.Foot_Ball_Team(Name)
```

- Optimization using query re-writing:

```
SELECT COUNT(*)
FROM dbo.Plays
INNER JOIN dbo.Player ON Player.PlayerID = Plays.PlayerID
INNER JOIN dbo.Foot_Ball_Team ON Foot_Ball_Team.TeamID = Player.TeamID
WHERE dbo.Foot_Ball_Team.Name LIKE 'AB%'
```



- Execution plan:



- Theoretical parallel query processing:

The following can be executed in parallel:

- Index seek on Player table
- Index seek on Foot\_Ball\_Team table
- Clustered Index seek on Plays table

# Optimization Details:

The new database statistics after modification:

- Disk Usage by Partition:

Table Name	# Records	Reserved (KB)	Used (KB)
[-] dbo.Foot_Ball_Team	1,000,000	104,344	75,144
[-] Index (PK__Foot_Bal__123AE7B93A9770BB)	1,000,000	65,160	36,088
[-] Index (ix_footballA)	1,000,000	17,928	17,896
[-] Index (ix_footteam_name)	1,000,000	21,256	21,160
[-] dbo.Game	1,000,000	79,696	70,056
[-] Index (PK__Game__2AB897DD900A37E4)	1,000,000	65,736	56,160
[-] Index (ix_MNTH)	1,000,000	13,960	13,896
[-] dbo.League	1,000,000	78,864	42,016
[-] Index (PK__League__10ABBC14D2017D91)	1,000,000	64,904	28,120
[-] Index (indF)	1,000,000	13,960	13,896
[-] dbo.Player	1,000,000	144,464	95,624
[-] Index (PK__Player__4A4E74A83ADD5CF7)	1,000,000	130,504	81,720
[-] Index (ix_player_TID)	1,000,000	13,960	13,904
[-] dbo.Plays	1,000,000	65,544	36,776
[-] Index (PK__Plays__88E5FDD575D92846)	1,000,000	65,544	36,776
[-] dbo.Plays_a	1,000,000	65,288	31,696
[-] Index (PK__Plays_a__D0916EC47299211C)	1,000,000	65,288	31,696
[-] dbo.Purchase_Tickets_Entrance	1,000,000	65,288	31,696
[-] Index (PK__Purchase__4999D91FE8CEFA4E)	1,000,000	65,288	31,696
[-] dbo.Round	1,000,000	64,712	20,872
[-] Index (PK__Round__94D84E1A6045E08A)	1,000,000	64,712	20,872
[-] dbo.Userr	1,000,000	65,672	54,520
[-] Index (PK__Userr__3214EC27F76887D0)	1,000,000	65,672	54,520

- Disk Usage by Table:

Table Name	# Records	Reserved (KB)	Data (KB)	Indexes (KB)	Unused (KB)
dbo.Foot_Ball_Team	1,000,000	104,344	35,944	39,200	29,200
dbo.Game	1,000,000	79,696	55,944	14,112	9,640
dbo.League	1,000,000	78,864	28,008	14,008	36,848
dbo.Player	1,000,000	144,464	81,408	14,216	48,840
dbo.Plays	1,000,000	65,544	36,592	184	28,768
dbo.Plays_a	1,000,000	65,288	31,536	160	33,592
dbo.Purchase_Tickets_Entrance	1,000,000	65,288	31,536	160	33,592
dbo.Round	1,000,000	64,712	20,784	88	43,840
dbo.Userr	1,000,000	65,672	54,312	208	11,152

# The enhancements made:

## The enhancements in the schema:

No enhancements were needed to be made in the schema

## The enhancement in memory management:

We used stored procedures because they can be used as a modular programming technique, which means create once, store, and then use several times whenever required. This means faster execution and no need to write and store different .sql file types that might get lost.

We decreased the size of varchar attributes from 55 to 20, which impacted the size of the database. Additionally, any query that tried to access any varchar typed attribute is now executed much faster.

## The modification in the indexes:

The following indices were made to enhance the query performance:

- `CREATE INDEX indf ON dbo.League(year)`
- `CREATE INDEX ix_footballA ON dbo.Foot_Ball_Team(Goals,LeagueID)`
- `CREATE INDEX ix_MNTH ON dbo.Game(Monthh)`
- `CREATE INDEX ix_footteam_name ON dbo.Foot_Ball_Team(Name)`
- `CREATE INDEX ix_player_TID ON dbo.Player(TeamID)`

## Modifications on the query statements:

For query 2, the inner query was removed and replaced with join conditions and a condition on the “Monthh” field of the Game table. Note that SQL server implicitly converts the join conditions to inner join.

For query 3, the scalar-valued function “SUBSTRING” was replaced with the “LIKE” operator. Index seek operators cannot be used in general with scalar-valued functions, which is why they should be avoided.

# Validation Details

## Effect of modifications on query performance:

Queries	Non-Optimized	Query Optimization	Indexing	Stored Procedure	Reduce col. size
Query 1	CPU: 3782	No query optimization	CPU: 3297	CPU: 2987	CPU: 2672
	Elapsed: 6694		Elapsed: 5677 [15.2%]	Elapsed: 5449[18.6%]	Elapsed: 5091[23.9%]
	Memory:88 Kb		Memory:96 Kb [-9%]	Memory:96 Kb [-9%]	Memory:96 Kb [-9%]
Query 2	CPU: 3453	CPU: 2312	CPU: 2110	CPU: 1937	CPU: 1828
	Elapsed: 9362	Elapsed: 7781[16.8%]	Elapsed: 7665[18.1%]	Elapsed: 7173[23.38%]	Elapsed: 6952[25.74%]
	Memory:64 Kb	Memory:48 Kb [25%]	Memory:48 Kb [25%]	Memory:48 Kb [25%]	Memory:48 Kb [25%]
Query 3	CPU: 1766	CPU: 1703	CPU: 1579	CPU: 1141	CPU: 1000
	Elapsed: 2759	Elapsed: 1946[29.47%]	Elapsed: 1753[36.46%]	Elapsed: 1133[58.9%]	Elapsed: 1012[63.3%]
	Memory:56 Kb	Memory:40 Kb [29%]	Memory:40 Kb [29%]	Memory:40 Kb [29%]	Memory:40 Kb [29%]

(Time is measured in ms)


## Effect of database size on query performance:

Queries	10,000	100,000	1,000,000	10,000,000
Query 1	CPU: 110	CPU: 343	CPU: 2672	CPU: 34812
	Elapsed: 226	Elapsed: 687	Elapsed: 5091	Elapsed: 58147
	Memory:96 Kb	Memory:96 Kb	Memory:96 Kb	Memory:96 Kb
Query 2	CPU: 47	CPU: 234	CPU: 1828	CPU: 26296
	Elapsed: 253	Elapsed: 1019	Elapsed: 6952	Elapsed: 73637
	Memory:64 Kb	Memory:48 Kb	Memory:48 Kb	Memory:48 Kb
Query 3	CPU: 32	CPU: 141	CPU: 1000	CPU: 11180
	Elapsed: 93	Elapsed: 231	Elapsed: 1012	Elapsed: 13506
	Memory:40 Kb	Memory:40 Kb	Memory:40 Kb	Memory:40 Kb

(Time is measured in ms)

## SQL vs NoSQL:

Comparisons were made on a database of size 1,000,000 rows, however, for query 1, we had to use a database of size 100,000 because we got an error converting to BSON as the size of the object exceeded the 16 MBs limit.



## Equivalent NoSQL queries used:

### Query1:

```
db.dbo.Foot_Ball_Team.find({
  Goals: {$in:
    db.dbo.Foot_Ball_Team.find({
      LeagueID: {$in:
        db.dbo.League.find({
          year: {$gt: 2007}
        }).map(l => l.LeagueID)
      },
      Goals: {$gt: 15000000}
    }).map(t=>t.Goals)
  },
  TeamID: {$in:
    db.dbo.Plays_a.find({
      GameID: {$in:
        db.dbo.Game.find().map(g => g.GameID)
      }
    }).map(p => p.TeamID)
  }
})
```

```
}).map(x => x.Name).length
```

## Query 2:

```
db.dbo.Foot_Ball_Team.find({
  TeamID: {$in:
    db.dbo.Plays_a.find({
      GameID: {$in:
        db.dbo.Game.find({
          Monthh: {$gt: 5}
        }).map(g => g.GameID)
      }
    }).map(p => p.TeamID)
  }
}).map(t => t.Name)
```

## Query 3:

```
db.dbo.Plays.find({
  PlayerID: {$in:
    db.dbo.Player.find({
      TeamID: {$in:
        db.dbo.Foot_Ball_Team.find({
          Name: /^AB.*/i
        }).map(t => t.TeamID)
      }
    }).map(p => p.PlayerID)
  }
}).map(x=>x).length
```

## Comments and Recommendations:

- We would recommend using SQL server as it does some pretty good optimizations to the queries implicitly, and it also suggests generating indices when applicable to enhance the query performance.
- We would also recommend trying to understand the execution plan and try to look for any index scans being performed and perhaps generating a non-clustered index to convert the scanning operations into seeking operations, which would enhance the query performance.