



Faculty Of Engineering



Cairo University

# Advanced Database System Report

Ahmed Mohamed Ismail Nabeel

1180501

Mostafa Ashraf Ahmed

1180406

Moaz Mohamed Elsherbini

1180528

Nader Youhanna Adib

1180477

## Table of Contents

System Information .....	4
Query 1.....	5
Before Optimization.....	5
Execution Plan.....	5
After Optimization .....	5
Execution Plan.....	6
Optimization Done.....	6
Theoretical Parallel Query Processing Report .....	6
Query 2.....	7
Before Optimization.....	7
Execution Plan.....	7
After Optimization .....	7
Execution plan.....	8
Optimization Done.....	8
Theoretical Parallel Query Processing Report .....	8
Query 3.....	9
Before Optimization.....	9
Execution Plan.....	9
After Optimization .....	9
Stored Procedure .....	10
Execution Plan.....	10
Optimization Done.....	10
Theoretical Parallel Query Processing Report .....	10
Database Statistics .....	11
Times for Different Database Sizes.....	12
Time Analysis .....	13
NoSQL Queries .....	15
Query 1.....	15
Query 2.....	16
Query 3.....	18
Optimizations.....	19
Schema Enhancement .....	19
Memory and Cache Management Enhancement.....	19

Indexes Modifications.....	19
Query Rewriting Modifications.....	19
Conclusion.....	19

## System Information

**Processor** Intel(R) Core(TM) i7-3612QM CPU @ 2.10GHz 2.10 GHz

**Installed RAM** 16.0 GB (15.9 GB usable)

Operating System: Windows 10 Pro

Hard disk: 1 TB SSD

Item	Value
OS Name	Microsoft Windows 10 Pro
Version	10.0.19045 Build 19045
Other OS Description	Not Available
OS Manufacturer	Microsoft Corporation
System Name	AHMED-DELL
System Manufacturer	Dell Inc.
System Model	Inspiron 7520
System Type	x64-based PC
System SKU	Inspiron 7520
Processor	Intel(R) Core(TM) i7-3612QM CPU @ 2.10GHz, 2101 Mhz, 4 Core(s), 8 Logical ...
BIOS Version/Date	Dell Inc. A14, 17-May-18
SMBIOS Version	2.7
Embedded Controller Version	1.01
BIOS Mode	Legacy
BaseBoard Manufacturer	Dell Inc.
BaseBoard Product	0PXH02
BaseBoard Version	A00
Platform Role	Mobile
Secure Boot State	Unsupported
PCR7 Configuration	Binding Not Possible
Windows Directory	C:\Windows
System Directory	C:\Windows\system32
Boot Device	\Device\HarddiskVolume1
Locale	United States
Hardware Abstraction Layer	Version = "10.0.19041.2251"

**Installed Physical Memory (RAM)** 16.0 GB

**Total Physical Memory** 15.9 GB

**Available Physical Memory** 6.59 GB

**Total Virtual Memory** 18.3 GB

**Available Virtual Memory** 7.63 GB

**Page File Space** 2.38 GB

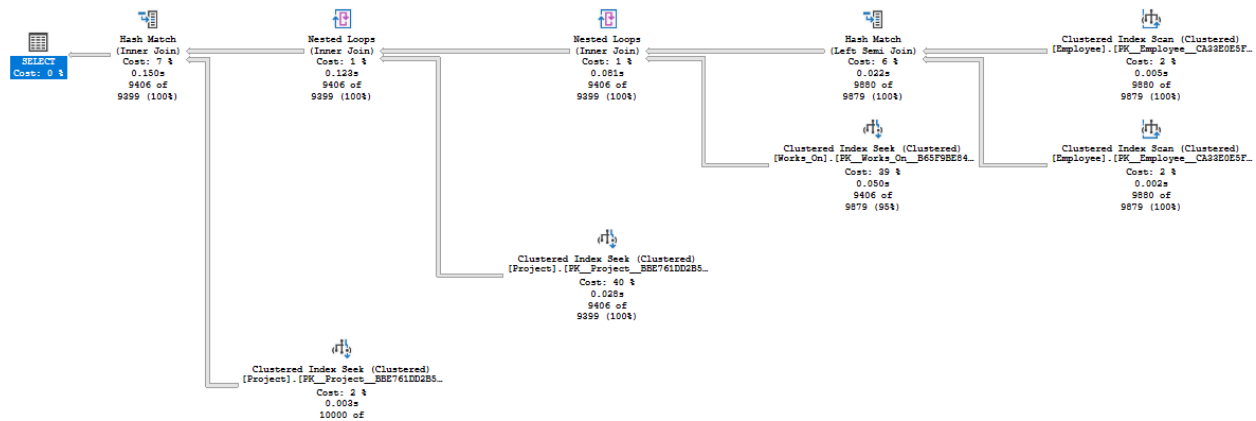
**Page File** F:\pagefile.sys

**Kernel DMA Protection** Off

## Before Optimization

```
SELECT E.Fname, E.Salary, E.Super_Ssn, P.Pname, W.Hours
FROM Employee E INNER LOOP JOIN Works_On W ON E.Ssn = W.Essn
                INNER LOOP JOIN Project P ON W.Pno = P.Pnumber
WHERE W.Hours > 500
      AND E.Salary in (
        SELECT E.Salary FROM Employee E
        WHERE E.Salary > 10000
      )
      AND P.Pnumber in (
        SELECT P.Pnumber FROM Project P
        WHERE P.Pnumber > 500
      )
GO
```

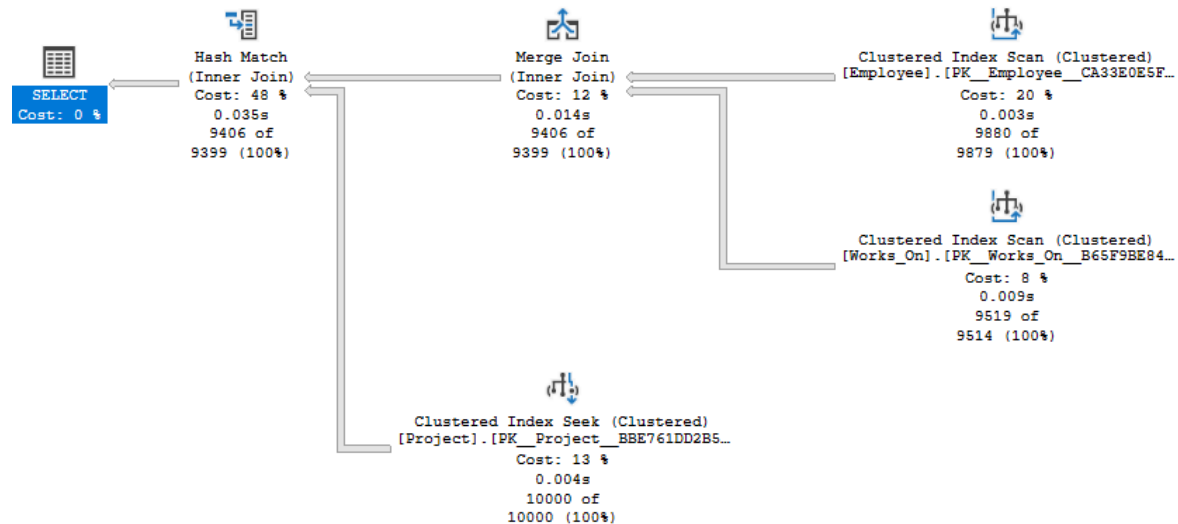
## Execution Plan



## After Optimization

```
SELECT E.Fname, E.Salary, E.Super_Ssn, P.Pname, W.Hours
FROM
    Employee E
    INNER JOIN Works_On W ON E.Ssn = W.Essn
    INNER JOIN Project P ON W.Pno = P.Pnumber
WHERE
    E.Salary > 10000
    AND P.Pnumber > 500
    AND W.Hours > 500
GO
```

## Execution Plan



## Optimization Done

- INNER JOIN instead of NESTED LOOP JOIN
- Non-clustered index on Ssn & Pnumber

Action	CPU Time (s)	Elapsed Time (s)
Query 1 (No optimization)	0.688	2.180
Query 1 (Optimized INNER JOIN)	0.136	2.106
Adding non-clustering index	0.123	1.902
NoSQL MongoDB	2.223	

## Theoretical Parallel Query Processing Report

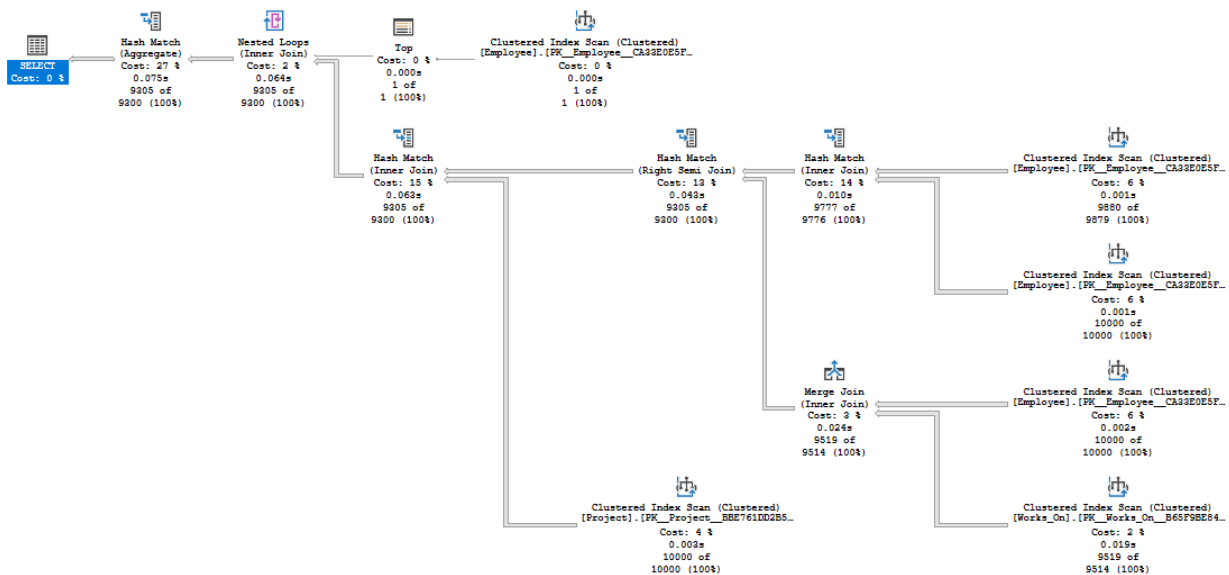
- Clustered index scan on Employee table
- Clustered index scan on Project table
- Clustered index scan on Works\_On table

## Query 2

### Before Optimization

```
SELECT DISTINCT E.Fname, E.Salary, E.Super_Ssn, P.Pname, W.Hours
FROM Employee E, Employee M, Works_On W, Project P
WHERE
    W.Hours > 500
    AND E.Ssn = W.Essn
    AND W.Pno = P.Pnumber
    AND E.Super_Ssn in (
        SELECT
            M.Ssn
        FROM
            Employee E,
            Employee M
        WHERE
            E.Super_Ssn = M.Ssn
            AND E.Salary > 10000
    )
GO
```

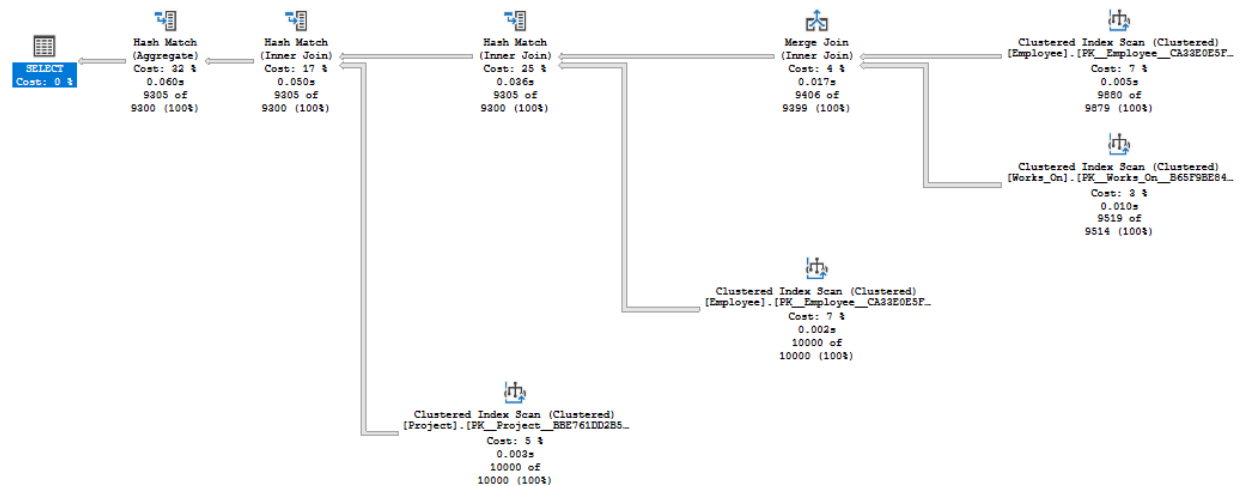
### Execution Plan



### After Optimization

```
SELECT DISTINCT E.Fname, E.Salary, E.Super_Ssn, P.Pname, W.Hours
FROM Employee E
    INNER join Employee M on E.Super_Ssn = M.Ssn
    inner join Works_On W on E.Ssn = W.Essn
    INNER join Project P on W.Pno = P.Pnumber
WHERE
    Hours > 500
    AND E.Salary > 10000
GO
```

## Execution plan



## Optimization Done

- INNER JOIN instead of nested loop.
- Non-clustered index on Ssn, Pnumber, Salary and Hours.

Action	CPU Time (s)	Elapsed Time (s)
Query 1 (No optimization)	0.476	4.467
Query 1 (Optimized INNER JOIN)	0.438	1.922
Adding non-clustering index	0.317	1.800
NoSQL MongoDB	310.751	

## Theoretical Parallel Query Processing Report

- Clustered index scan on Project table
- Clustered index scan on Works\_On table
- Clustered index scan on Employee

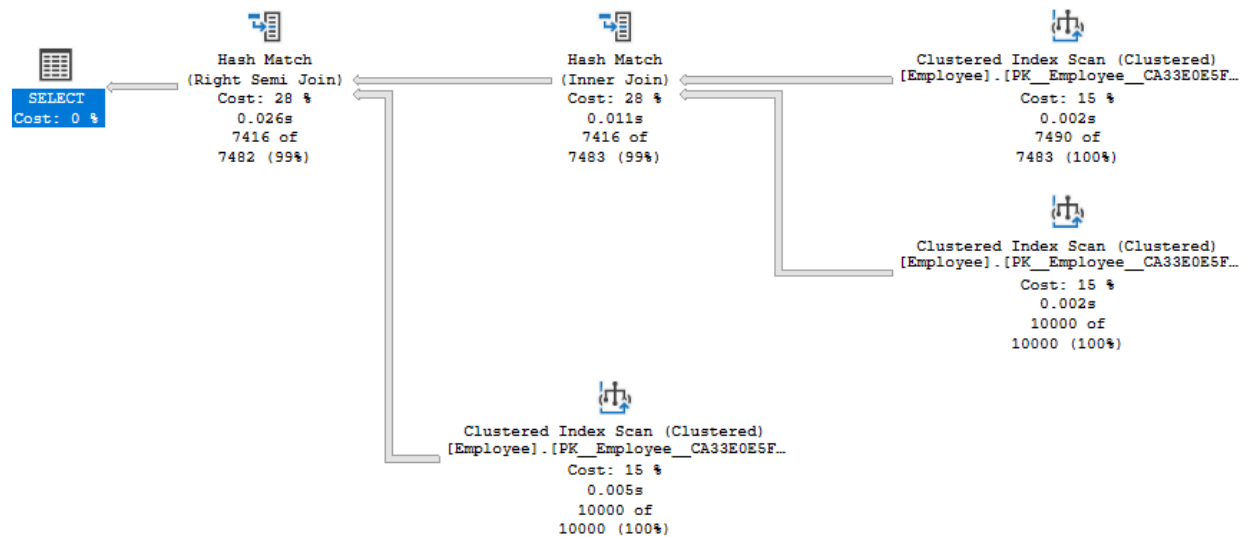


## Query 3

### Before Optimization

```
SELECT * FROM Employee E
WHERE
    E.Super_Ssn IN (
        SELECT
            M.Ssn
        FROM
            Employee E,
            Employee M
        WHERE
            E.Super_Ssn = M.Ssn
            AND M.Bdate > '1970-01-01'
    );
```

### Execution Plan



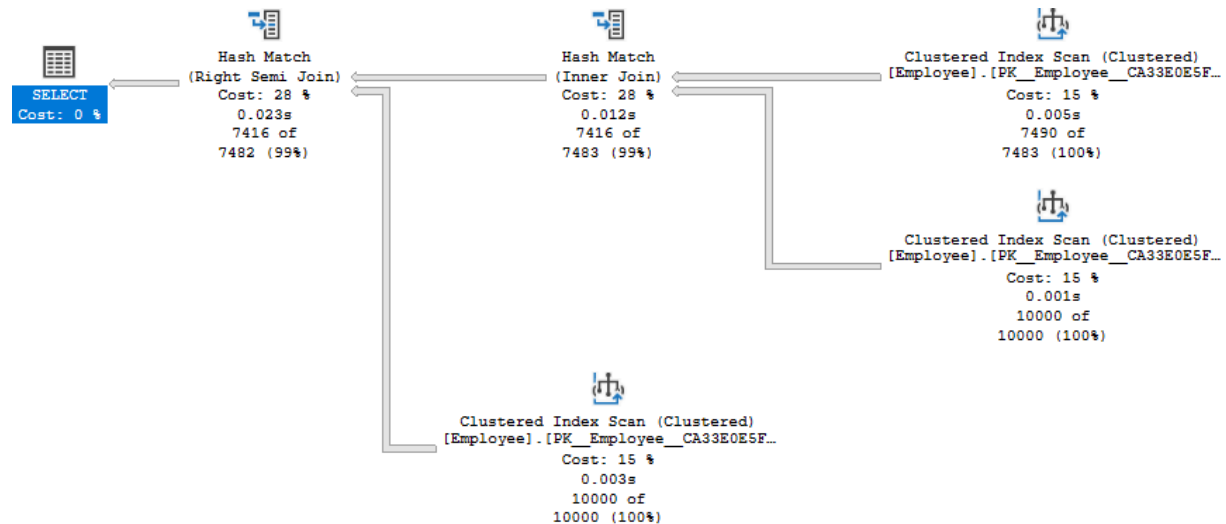
### After Optimization

```
EXEC sp_Employee_Super_Ssn '1970-01-01'
```

## Stored Procedure

```
CREATE PROCEDURE sp_Employee_Super_Ssn @Bdate DATETIME AS BEGIN
Select
    *
from
    Employee E
Where
    E.Super_Ssn in (
        Select
            M.Ssn
        From
            Employee E,
            Employee M
        where
            E.Super_Ssn = M.Ssn
            And M.Bdate > @Bdate
    )
END
```

## Execution Plan



## Optimization Done

- Used stored procedure.
- Non-clustered index on Ssn & Bdate.

Action	CPU Time (s)	Elapsed Time (s)
Query 1 (No optimization)	0.139	1.983
Query 1 (Stored Procedure)	0.145	1.751
Adding non-clustering index	0.121	2.002
NoSQL MongoDB	17.946	

## Theoretical Parallel Query Processing Report

- Clustered index scan on Employee



## Database Statistics


Table Name	Row Count	Main Key	Indexes	FK	Identity Column	Max Row Size (Bytes)
Employee	1M	Yes	1	2	Yes	71
Department	1M	Yes	1	1	Yes	19
Project	1M	Yes	1	1	Yes	38
Works_On	1M	No	1	2	No	12
Department_Location	1M	No	1	1	No	19
Dependent	1M	No	1	1	No	38

## Times for Different Database Sizes

Query Size	10K (s)	100K (s)	1M (s)	10M (s)
<b>Query 1 (Non-optimized)</b> CPU Time	0.688	2.183	19.771	184.889
<b>Query 1 (Non-optimized)</b> Elapsed Time	2.180	4.666	33.409	290.898
<b>Query 1 (Optimized)</b> CPU Time	0.123	0.769	7.126	62.678
<b>Query 1 (Optimized)</b> Elapsed Time	1.902	3.315	20.760	170.614
<b>Query 2 (Non-optimized)</b> CPU Time	0.476	1.739	12.106	120.700
<b>Query 2 (Non-optimized)</b> Elapsed Time	4.467	4.348	25.066	230.224
<b>Query 2 (Optimized)</b> CPU Time	0.317	0.997	10.092	101.365
<b>Query 2 (Optimized)</b> Elapsed Time	1.800	3.495	22.735	209.833
<b>Query 3 (Non-optimized)</b> CPU Time	0.139	0.944	7.204	76.916
<b>Query 3 (Non-optimized)</b> Elapsed Time	1.983	3.432	19.330	187.511
<b>Query 3 (Optimized)</b> CPU Time	0.121	0.976	7.184	73.292
<b>Query 3 (Optimized)</b> Elapsed Time	2.002	3.542	19.047	179.315

## Time Analysis





## NoSQL Queries

### Query 1

```
db.getCollection("Employee").aggregate([
  {
    "$project" : {
      "_id" : NumberInt(0),
      "E" : "$$ROOT"
    }
  },
  {
    "$lookup" : {
      "localField" : "E.Ssn",
      "from" : "Works_On",
      "foreignField" : "Essn",
      "as" : "W"
    }
  },
  {
    "$unwind" : {
      "path" : "$W",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$lookup" : {
      "localField" : "W.Pno",
      "from" : "Project",
      "foreignField" : "Pnumber",
      "as" : "P"
    }
  },
  {
    "$unwind" : {
      "path" : "$P",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$match" : {
      "E.Salary" : {
        "$gt" : NumberLong(10000)
      },
      "P.Pnumber" : {
        "$gt" : NumberLong(500)
      },
      "W.Hours" : {
        "$gt" : NumberLong(500)
      }
    }
  },
  {
    "$project" : {
      "E.Fname" : "$E.Fname",
      "E.Salary" : "$E.Salary",
      "E.Super_Ssn" : "$E.Super_Ssn",
      "P.Pname" : "$P.Pname",
      "W.Hours" : "$W.Hours",
      "_id" : NumberInt(0)
    }
  }
],
{
  "allowDiskUse" : true
});
```

## Query 2

```
db.getCollection("Employee").aggregate([
  {
    "$project" : {
      "_id" : NumberInt(0),
      "E" : "$$ROOT"
    }
  },
  {
    "$lookup" : {
      "localField" : "E.Super_Ssn",
      "from" : "Employee",
      "foreignField" : "Ssn",
      "as" : "M"
    }
  },
  {
    "$unwind" : {
      "path" : "$M",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$lookup" : {
      "localField" : "E.Ssn",
      "from" : "Works_On",
      "foreignField" : "Essn",
      "as" : "W"
    }
  },
  {
    "$unwind" : {
      "path" : "$W",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$lookup" : {
      "localField" : "W.Pno",
      "from" : "Project",
      "foreignField" : "Pnumber",
      "as" : "P"
    }
  },
  {
    "$unwind" : {
      "path" : "$P",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$match" : {
      "W.Hours" : {
        "$gt" : NumberLong(500)
      }
    }
  }
])
```



```

        "E.Salary" : {
            "$gt" : NumberLong(10000)
        }
    },
    {
        "$project" : {
            "E.Fname" : "$E.Fname",
            "E.Salary" : "$E.Salary",
            "E.Super_Ssn" : "$E.Super_Ssn",
            "P.Pname" : "$P.Pname",
            "W.Hours" : "$W.Hours",
            "_id" : NumberInt(0)
        }
    },
    {
        "$group" : {
            "_id" : null,
            "distinct" : {
                "$addToSet" : "$$ROOT"
            }
        }
    },
    {
        "$unwind" : {
            "path" : "$distinct",
            "preserveNullAndEmptyArrays" : false
        }
    },
    {
        "$replaceRoot" : {
            "newRoot" : "$distinct"
        }
    }
],
{
    "allowDiskUse" : true
}
);

```

### Query 3

```
db.getCollection("Employee").aggregate([
  {
    "$project" : {
      "_id" : NumberInt(0),
      "E" : "$$ROOT"
    }
  },
  {
    "$lookup" : {
      "localField" : "E.non_existing_field",
      "from" : "Employee",
      "foreignField" : "non_existing_field",
      "as" : "M"
    }
  },
  {
    "$unwind" : {
      "path" : "$M",
      "preserveNullAndEmptyArrays" : false
    }
  },
  {
    "$match" : {
      "$and" : [
        {
          "$expr" : {
            "$eq" : [
              "$E.Super_Ssn",
              "$M.Ssn"
            ]
          }
        },
        {
          "$expr" : {
            "$gt" : [
              "$M.Bdate",
              "1970-01-01"
            ]
          }
        }
      ]
    }
  },
  {
    "allowDiskUse" : true
  }
]);
```

## Optimizations

### Schema Enhancement

- We used the same schema used in Phase 1

### Memory and Cache Management Enhancement

- We used stored procedures to enhance memory and cache management

### Indexes Modifications

- We added some non-clustered indexes to speed up the data selection as follows:
  - Non-clustered index on: Ssn (Employee Table)
  - Non-clustered index on: Pnumber (Project Table)
  - Non-clustered index on: Bdate (Works\_on Table)
  - Non-clustered index on: Salary (Employee Table)

### Query Rewriting Modifications

- We rewrote the queries to replace the nested loops with INNER JOIN to make the query execute faster

## Conclusion

- Using SQL server is recommended for relational schemas as NoSQL has worse performance
- Index tuning has a great effect in the execution time for executing the queries
- Optimization for SQL query has a more noticeable effect for larger database sizes
- Stored procedures did not yield the expected optimization
- Seeing the execution plan and the index scan is very useful when trying to generate non-clustered indexes to convert the scanning operations into seeking operations.